

# Concurrent and Parallel Programming, Part I

Programming Languages  
CS 214

(1/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Introduction

Suppose you and that “special someone” are shopping for champagne and oysters for a romantic dinner...

You might do your shopping either of two ways:

- The “we can’t bear to be apart” approach: together, you
  - find the oysters,
  - find the champagne,
  - pay for what you’ve selected at the checkout
- The “unromantic but efficient” divide-and-conquer approach:
  - one of you finds the oysters,
  - the other finds the champagne,
  - rendezvous at the checkout to pay for what you’ve selected

If the average time to find champagne+oysters sequentially is  $\tau$ , then finding them *concurrently* takes about \_\_\_\_.

(2/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Concurrent Programming

Most modern programming languages provide built-in support for concurrent processing.

- \_\_\_\_\_ provides a *Task* construct, and each distinct task is automatically executed concurrently.
- \_\_\_\_\_ provides a *Thread* class, and each distinct thread can be executed concurrently.
- \_\_\_\_\_ added a standard *thread* class in C++11.
- \_\_\_\_\_, ... provide multithreading/processing capabilities

Older languages:

- \_\_\_\_\_ relies on external libraries for concurrency (e.g., Unix *fork()*, POSIX *pthreads*, OpenMP, MPI, ...).
- \_\_\_\_\_ provides a *start-process* function, but it is *not standard Lisp*.

(3/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Example: Ada

We might represent our sequential approach as follows:

```
procedure RomanticApproach is
begin
  FindOysters; ← _____ do this
  FindChampagne; ← _____
  PayAtCheckout; ← _____
end RomanticApproach;
```

By contrast, our concurrent approach is:

```
procedure EfficientApproach is
  task OysterFinder;
  task body OysterFinder is begin
    FindOysters; RendezvousAtCheckout; } do this
  end OysterFinder;
begin
  FindChampagne; RendezvousAtCheckout; } _____
end EfficientApproach;
```

(4/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Example: Java

In Java, our sequential approach is expressed as:

```
class RomanticApproach {  
    public static void main(String [] args) {  
        findOysters();  
        findChampagne();  
        payAtCheckout();  
    }  
}
```

By contrast, our concurrent approach is expressed as:

```
class EfficientApproach {  
    class OysterFinder extends Thread {  
        public void run() {  
            findOysters(); rendezvousAtCheckout();  
        }  
    }  
    public static void main(String [] args) {  
        OysterFinder of = new OysterFinder();  
        of.start();  
        findChampagne(); rendezvousAtCheckout();  
    }  
}
```

(5/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Terminology

- A \_\_\_\_\_ is a computer with one processing core...
  - With a time-sharing OS, concurrent processing results in \_\_\_\_\_ (aka *logical concurrency*), because the OS time-shares the single core among the processes/threads.
- A \_\_\_\_\_ is a computer with multiple cores...
  - Concurrent processing results in \_\_\_\_\_ (aka *true concurrency*), as the OS can simultaneously run different processes/threads on different cores.
  - In a \_\_\_\_\_ *multiprocessor*, the cores
    - \_\_\_\_\_, and
    - are usually in *close physical proximity*.
  - In a \_\_\_\_\_ *multiprocessor*, the cores
    - have \_\_\_\_\_ (each has its own *local memory*),
    - are often *not in close physical proximity*.

(6/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College




## Parallelizing Compilers

What *dependencies* exist in the following function?

```
void f(double x) {
    double y, z;    // 1
    fin >> y;        // 2 ... depends on 1
    z = x * x;        // 3 ... depends on 1
    return y * z;     // 4 ... depends on 2, 3
}
```

Parallelizing compilers build \_\_\_\_\_,  
and use them to identify pieces of code that can be  
executed concurrently:

- branches in the graph (e.g., 2 & 3)  
can be safely executed in parallel
- CPU must have the extra hardware to  
perform \_\_\_\_\_



(7/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Processes and Threads

A computation is sometimes  
called a \_\_\_\_\_:

```
int main() {
    s1;
    s2;
    ...
    sn;
}
```

The sequence of *events* that  
occur as control flows  
through a process is called a  
\_\_\_\_\_:

```
s1; }
s2; }
... }
sn; }
```

If, for the same inputs, the sequence of events always occurs  
in the same order, the sequence is called \_\_\_\_\_;  
otherwise, the sequence is called \_\_\_\_\_.

(8/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Goal: Speedup

One goal of concurrent processing is to achieve speedup:

If a task requires  $\tau$  time-units to solve sequentially,  
but can be split into  $p$  subtasks that can be solved in parallel,  
then parallel processing can perform it in \_\_\_\_\_ time-units.

Formally, speedup can be define as:

\_\_\_\_\_

where:  $T_1$  is the time 1 task takes to solve the problem, and  
 $T_N$  is the time  $N$  tasks take to solve the problem.

If it takes one person 10 minutes to find champaign+oysters,  
but it takes two people 6 minutes, the speedup is \_\_\_\_\_.

(9/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Goal: Responsiveness

... is another goal of concurrent processing.

Example 1: \_\_\_\_\_

- If an application has a single thread and it has to perform a time-consuming task, then the GUI will “freeze” while the thread is performing that task.
- If the application uses one thread to handle user-interface events and forks a separate thread to perform each task, then the GUI will remain responsive.

Example 2: \_\_\_\_\_

- If a server has a single thread and has to perform a time-consuming task, then the server will be unable to accept incoming requests while it is performing the task.
- If the server uses one thread to accept incoming requests and forks a new thread to handle each request, the server will accept and handle all requests it receives.

(10/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Event Interleaving

When a computation consists of two or more processes:

```
void p() {
  a;
  b;
  c;
}
```











```
void q() {
  x;
  y;
}
```

then the events in the threads of those processes may be \_\_\_\_\_:

– Each sequence is

\_\_\_\_\_,  
but the *sequencing across both processes* is

\_\_\_\_\_.



(11/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Interleaving in Ada

We can see non-deterministic behavior in Ada as follows:

```
procedure Interleave is
  task A;
  task body A is begin
    put("a");
  end A;
  task B;
  task body B is begin
    put("b");
  end B;
  task C;
  task body C is begin
    put("c");
  end C;
begin
  null;
end Interleave;
```

Sample Run:

```
% ./interleave
% ./interleave
% ./interleave
% ./interleave
...
```

How many distinct executions are there?

→ The number of \_\_\_\_\_ of {a, b, c}

→ Discrete Math: a set of  $n$  elements has \_\_\_\_\_  
permutations, so \_\_\_\_\_ distinct possibilities.

(12/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science


Calvin College



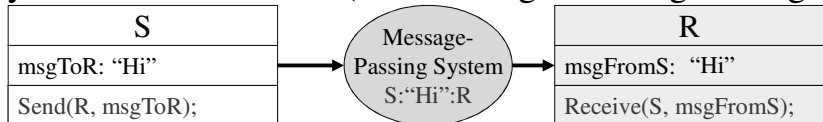
## Communication

If a language allows us to divide a task into subtasks, it should also provide a way for those tasks to *communicate*.

- On a \_\_\_\_\_ multiprocessor, two tasks can communicate through the \_\_\_\_\_:



- On a \_\_\_\_\_ multiprocessor, they can communicate by \_\_\_\_\_ (i.e., sending-receiving messages):



(13/39)

©Joel C. Adams. All Rights Reserved.


Dept of Computer Science

Calvin College




## Shared Memory Synchronization

If two tightly-coupled processes try to access shared memory simultaneously, the result may be incorrect... What can go wrong?



Accesses to shared memory must be \_\_\_\_\_ to avoid this:



Synchronization forces one process to \_\_\_\_\_ until the other is finished.

(14/39)

©Joel C. Adams. All Rights Reserved.


Dept of Computer Science

Calvin College




## Synchrony

Concurrent computations lie on a continuum, depending on how much communication/synchronization they entail:



- *Lock-step synchronous* computations must communicate or be re-synchronized \_\_\_\_\_ of the computation:



- \_\_\_\_\_ (aka *embarrassingly parallel*) computations require no communication/synchronization of their processes

(15/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College




## The \_\_\_\_\_ Problem

There are a number of “classic” synchronization problems, one of which is the producer-consumer problem...

There are  $M$  *producers* that put items into a buffer. The buffer is shared with  $N$  *consumers* that remove items from the buffer.

- The problem is to devise a solution that ensures \_\_\_\_\_.



Accesses to the buffer must be \_\_\_\_\_: if multiple producers / consumers access it simultaneously, producers may overwrite each other's values, consumers may retrieve the same value, etc.

In the *bounded-buffer version*, the buffer has a \_\_\_\_\_.

(16/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College






## The \_\_\_\_\_ Problem

... is another “classic” synchronization problem:

Five philosophers sit at a table, alternating between eating noodles and thinking. In order to eat, a philosopher must have two chopsticks. However, there is a single chopstick between each pair of plates, so if one is eating, neither neighbor can eat. A philosopher puts down both chopsticks when thinking.

–Devise a solution that ensures:

- no philosopher starves; and
- a hungry philosopher is only prevented from eating by his immediate neighbor(s).



```
task philosopher is
  while True begin
    think(randomTime);
    get(left);
    get(right);
    eat();
    release(left);
    release(right);
  end while;
End philosopher;
```

(17/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science


Calvin College



## The Dining Philosophers (ii)

How about this instead?

```
task philosopher is begin
  while True begin
    think(randomTime);
    while not (have(left) and have(right))
      begin
        get(left);
        if notInUse(right) then
          get(right);
        else
          release(left);
        end if;
      end while;
    eat;
    release(left);
    release(right);
  end while;
end philosopher;
```



(18/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Mutual Exclusion

An object that can only be accessed by “one-thing-at-a-time” (e.g., shared memory) is called a \_\_\_\_\_ object:

- An access by one process *excludes* other processes from access
- A task may have to ‘wait its turn’ to access the object
- Many real-world objects (e.g., chopsticks) are mutually exclusive
- Shared-memory *writes* are mutually exclusive: \_\_\_\_\_!

TravelAgent1		TravelAgent2
if (emptySeats > 0) emptySeats--; else display(“no more seats!”); end if;	emptySeats: 1	if (emptySeats > 0) emptySeats--; else display(“no more seats!”); end if;

- Shared-memory *reads* are not mutually exclusive, unless any task tries to write to the shared memory: \_\_\_\_\_!

(19/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Synchronization Primitives

In 1965, Dijkstra proposed the \_\_\_\_\_: a shared-memory programming mechanism that can be used to synchronize accesses to a mutually-exclusive resource, with two values: {*locked*, *unlocked*}, and three simple operations:

- *Initialize* the semaphore to *unlocked*
- P: *Lock* the semaphore (wait if it is already *locked*)
- V: *Unlock* the semaphore (awaken the first process waiting for it).

Java 1.7 added a *Semaphore* class:

Operation	Dijkstra	Java Syntax
Initialization (unlocked)	s: Semaphore;	s = new Semaphore(1);
Lock the semaphore		
Unlock the semaphore		

(20/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Semaphores and Mutual Exclusion

A semaphore is a shared-memory variable that can be used to \_\_\_\_\_  
\_\_\_\_\_ to other shared-memory variables:

TravelAgent1		TravelAgent2
_____	emptySeats: 1	_____
if (emptySeats > 0)	s: unlocked	if (emptySeats > 0)
emptySeats--;		emptySeats--;
else		else
display("no more seats!");		display("no more seats!");
_____		_____

Whichever travel agent executes P(s) *first* (even by a nanosecond) will lock s, decrement emptySeats, and then unlock s.

Whichever travel agent executes P(s) *second* will find s locked and have to wait until the other agent unlocks s, discover that emptySeats == 0, and then get the "no more seats" message.

(21/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Semaphores and Lockstep Synchrony

A semaphore also permits two processes to execute in lock-step:

S		R
msgToR:	sharedSpace:	msgFromS:
loop {	okToWrite: unlocked	loop {
_____	okToRead: locked	_____
sharedSpace = msgToR;		msgFromS = sharedSpace;
_____		_____
}		}

- If R executes first, it will wait on the (*locked*) *okToRead* semaphore, until S signals (after it writes a value to the shared memory)...
- If S executes first, it will lock the (*unlocked*) *okToWrite* semaphore, write its message to shared memory, signal *okToRead*, and then wait on the (*locked*) *okToWrite* until R signals (after it finishes reading).

A \_\_\_\_\_ is a group of statements that access shared space.

(22/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science


Calvin College



## Dining Philosophers (iii)

What if we use a Semaphore?

```
Semaphore: wantBothSticks;  
task philosopher is begin  
  while True begin  
    think(randomTime);  
    while not haveBothSticks begin  
      P(wantBothSticks);  
      if available(leftStick) and  
        available(rightStick) then begin  
        get(left);  
        get(right);  
      end if;  
      V(wantBothSticks);  
    end while;  
    eat();  
    release(leftStick);  
    release(rightStick);  
  end while;  
end philosopher;
```



This seems to do  
it, but synchrony  
is so tricky, it's  
hard to be 100%  
certain it's correct  
...

(23/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Locks and Condition Variables

A semaphore may be used for either of two purposes:

- \_\_\_\_\_: guarding access to a critical section
- \_\_\_\_\_: making threads/processes suspend/resume

This dual use can lead to confusion: it may be unclear which role a semaphore is playing in a given computation...

For this reason, newer languages provide *distinct constructs for each*:

- \_\_\_\_\_: guarding access to a critical section
- \_\_\_\_\_: making threads wait until a condition is true

Locks support mutually-exclusive access to shared memory;  
condition variables support thread/process synchronization.

(24/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Locks

Like a Semaphore, a lock has two associated operations:

- *acquire()* - try to lock the lock; if it is locked, go to sleep
- *release()* - unlock the lock; awaken a waiting thread (if any)

The *acquire()* is analogous to the Semaphore \_\_\_\_\_ operation;  
the *release()* is analogous to the Semaphore \_\_\_\_\_ operation.

These can be used to 'guard' a critical section:

<pre>sharedLock.acquire(); // access sharedObj sharedLock.release();</pre>	<pre>Lock sharedLock; Object sharedObj;</pre>	<pre>sharedLock.acquire(); // access sharedObj sharedLock.release();</pre>
------------------------------------------------------------------------------------	---------------------------------------------------	------------------------------------------------------------------------------------

Every Java class inherits a *hidden lock* from class *Object*;  
the \_\_\_\_\_ keyword uses it:

```
synchronized {  
    // critical section  
}
```

(25/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Condition Variables

A *Condition* is a predefined type available in some languages that can be used to declare variables for synchronization.

When a thread needs to suspend execution inside a critical section until some condition is met, a *Condition* can be used.

There are three operations for a *Condition*:

- \_\_\_\_\_  
– suspend immediately; enter a queue of waiting threads
- \_\_\_\_\_, aka *notify()* in Java  
– awaken a waiting thread (usually the first in the queue), if any
- \_\_\_\_\_, aka *notifyAll()* in Java  
– awaken all waiting threads, if any

Every Java class inherits it from class *Object* a hidden condition-variable, and the *wait()*, *notify()* & *notifyAll()* methods that use it.

(26/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Monitors

Semaphores, Locks, and Conditions are simple but powerful synchronization tools; but many believe that they are *too powerful for the average programmer* (like the *goto*)...

— \_\_\_\_\_ are easy mistakes to make

Just as control structures were “higher level” than the *goto*, language designers began looking for higher level ways to synchronize processes.

In 1973, Brinch-Hansen and Hoare proposed the \_\_\_\_\_, a class whose methods are automatically accessed in a mutually-exclusive manner.

— A monitor *prevents simultaneous access by multiple threads*

(27/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Mesa-Style Monitors

*Concurrent Pascal* was first to provide a *Monitor* construct:

```
type BoundedBuffer = Monitor
  constant N := 1024;
  myHead, myTail, mySize: integer := 0;
  myValues : array(0..N-1) of Object;
  notEmpty, notFull: Condition;

  procedure put(obj: Object) begin
    while mySize = N do notFull.wait; end;
    myValues(myHead) := obj;
    myHead := (myHead + 1) mod N; mySize := mySize + 1;
    notEmpty.signal;
  end;

  procedure get(var obj: Object) begin
    while mySize = 0 do notEmpty.wait; end;
    obj = myValues(myTail);
    myTail = (myTail + 1) % N; mySize := mySize - 1;
    notFull.signal;
  end;
end;
```

(28/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College





## Buffer Synchrony

```
// ... continued from previous page ...
public synchronized void put(Object obj) {
    while ( this.isFull() )
        try{ wait(); } catch(Exception e) {}
    myValues[myHead] = obj;
    myHead = (myHead + 1) % myMax;
    mySize++;
    notifyAll();
}
public synchronized Object get() {
    Object result;
    while ( this.isEmpty() )
        try{ wait(); } catch(Exception e) {}
    result = myValues[myTail];
    myTail = (myTail + 1) % myMax;
    mySize--;
    notifyAll();
    return result;
}
}
```

The wait()  
operation  
causes the  
executing  
thread to

The  
notifyAll()  
operation

all waiting  
threads.

(31/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science


Calvin College



## Bounded Buffer Producer-Consumer

We can then use our *BoundedBuffer* as follows:

```
// producer thread
for (;;) {
    // produce Item it;
    buf.put(it);
}
```



```
// consumer thread
for (;;) {
    buf.get(it);
    // consume Item it;
}
```

\_\_\_\_\_ : No synchronization needed in producer or consumer.

Recall: Every Java class inherits a hidden *lock* from *Object*...

- When a *synchronized* method is called, it tries to *acquire* the lock (waiting if it is already locked), and *releases* the lock on termination.

Every Java class inherits a hidden *condition variable* from *Object*...

- wait() suspends a thread on the condition; notify() awakens a thread waiting on the condition; notifyAll() awakens all waiting threads.

(32/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College





## More Recent Java

In 2005, Java 1.5 added the package \_\_\_\_\_:

- A ThreadPool class and an Executor framework to make the management of groups of threads easier and more convenient
- Classes for thread-safe data structures (list, queue, map, ...)
- Classes for synchronization (semaphore, barrier, mutex, latch, ...)
- Classes for creating lock and condition variables;
- Classes for atomic operations (arithmetic, test-and-set, ...)

Subsequent Java releases have continued to add features:

- Futures, for asynchronous computations
- ForkJoinTasks and ForkJoinPools for recursive parallelism
- Lambda expressions, CompletableFutures, parallel streams, WorkStealingThreadPools for load-balancing, ...

(33/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## OpenMP ...

... stands for *Open MultiProcessing*

... is an industry-standard library for shared-memory parallel computing in C, C++, Fortran, ...

... uses \_\_\_\_\_

... simplifies the task of parallelizing legacy code

... was designed by a large consortium in 1997:  
*AMD, Cray, Fujitsu, HP, IBM, Intel, NEC, Nvidia, Oracle, Redhat, TI, ...*

... has “built in” support for many parallel design patterns

... continues to evolve (OpenMP 2.0 in 2000; 3.0 in 2008; 4.0 in 2014, ...; current version is 4.5)

(34/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Example: Summing an Array

```
#include <iostream>
#include <omp.h>

// getFileName(), readFile(), ...

int main(int argc, char** argv) {
    string filename = getFileName(argc, argv);
    vector<int> v = readFile(filename);

    #pragma omp parallel for reduction(+:sum)
    for (int i = 0; i < v.size(); i++) {
        sum += v[i];
    }

    cout << "The sum of the values in '"
         << filename << "' is " <<
         << sum << endl;
}
```

On a machine  
with  $t$  cores,  
this directive

\_\_\_\_\_ threads

An \_\_\_\_\_  
\_\_\_\_\_ occurs  
at the end of  
the stmt that  
follows the  
#pragma

(35/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College




## #pragma omp parallel for


The #pragma omp parallel directive forks the threads...

The for clause \_\_\_\_\_ that  
follows it across those threads.


On a single-core machine: \_\_\_\_\_ thread...



On a dual-core machine: \_\_\_\_\_ threads...



On a quad-core machine: \_\_\_\_\_ threads...



(36/39)

©Joel C. Adams. All Rights Reserved.

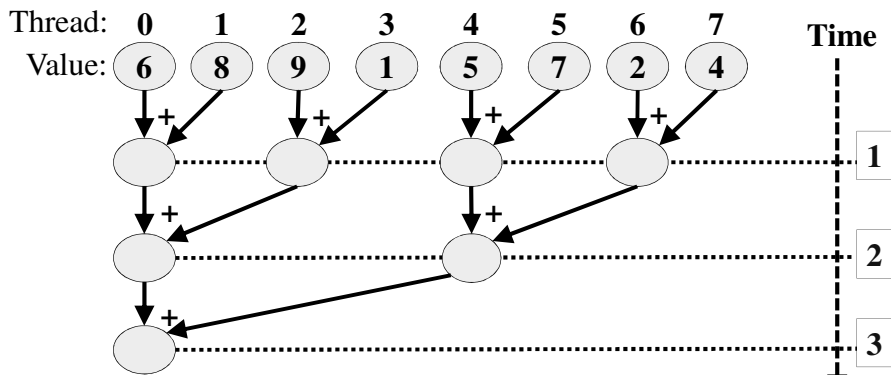
Dept of Computer Science

Calvin College



## The Reduction Clause

... in the `#pragma omp parallel for reduction(+:sum)` uses the `+` operator to combine the partial sums into variable `sum`.



Reduction reduces the sum-time from  $O(t)$  to \_\_\_\_\_ ...

(37/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## OpenMP also supports...

Other ways to reduce local results: `+`, `*`, `-`, `/`, `&`, `|`, `^`, `&&`, `||`, ...

Other directives: `#pragma omp` \_\_\_\_\_

- *critical*                      • *sections*                      • *task*                              • *single*
- *atomic*                        • *section*                        • *teams*                         • *master*
- *barrier*                      • *taskwait*                      • *simd*                         • ...

Library functions:

- *omp\_set\_num\_threads()*      • *omp\_init\_lock()*              • *omp\_get\_num\_teams()*
- *omp\_get\_num\_threads()*      • *omp\_set\_lock()*                • *omp\_get\_team\_size()*
- *omp\_get\_thread\_num()*        • *omp\_unset\_lock()*            • *omp\_get\_wtime()*
- *omp\_get\_num\_procs()*        • *omp\_test\_lock()*              • ...

Much, much more!

(38/39)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College

