

# Concurrent and Parallel Programming, Part II

Programming Languages  
CS 214

(1/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College




## Distributed Synchronization

Semaphores, locks, condition variables, monitors, are *shared-memory* constructs, and so \_\_\_\_\_.


– They are of \_\_\_\_\_ on a *distributed multiprocessor*

On a distributed multiprocessor, processes can communicate via \_\_\_\_\_ primitives.

– If the message-passing system has *no storage*, then the send/receive operations must be *synchronized*:



Two scenarios...



(2/23)

©Joel C. Adams. All Rights Reserved.


Dept of Computer Science

Calvin College




## Asynchronous Communication

–If the message-passing system has \_\_\_\_\_,  
then the send/receive operations can proceed \_\_\_\_\_:



The receiver can then retrieve the message when it is ready...

Two scenarios...



Message-buffering eliminates some (but not all) of the waiting.

(3/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science


Calvin College



## Message-Passing Languages

Some languages support message-passing between \_\_\_\_\_:

- \_\_\_\_\_ is a functional language developed at Ericsson and used by Nortel, T-Mobile, Facebook (chat, WhatsApp), and 20+ others.



```
B!{self(), {digits, [2,3,5]}}
```

```
receive
{A, {digits, nums}} ->
  analyze(nums);
end
```

- \_\_\_\_\_ is a hybrid OO+functional language used at Netflix, LinkedIn, Twitter, Tumblr, Foursquare, Sony, and other companies:

```
B! digits(2,3,5)
```

```
receive {
  case digits(nums) =>
    analyze(nums);
}
```

(4/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## An Ada Task

... has 3 characteristics:

- its own \_\_\_\_\_ of control;
- its own \_\_\_\_\_; and
- \_\_\_\_\_ (aka *entry procedures*)

Entry procedures are *self-synchronizing subprograms* that another task can invoke for task-to-task communication.

If task *t1* has an entry procedure *p*, then another task *t2* can execute:

*t1.p( argument-list );*

In order for *p* to execute, *t1* must execute:

*accept p ( parameter-list );*

- If *t1* executes *accept p* and *t2* has not called *p*, \_\_\_\_\_;
- If *t2* calls *p* and *t1* has not done *accept p*, \_\_\_\_\_.

(5/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science


Calvin College



## Rendezvous

When *t1* and *t2* are \_\_\_\_\_, *t1* *p* executes:

- *t2*'s *argument-list* is evaluated and passed to *t1.p*'s parameters
- *t2* suspends
- *t1* executes the body of *p*, using its parameter values
- return-values (or *out* or *in out* parameters) are passed back to *t2*
- *t1* continues execution;  
*t2* resumes execution



This interaction is called a \_\_\_\_\_ between *t1* and *t2*.

It does not depend on shared memory, so *t1* and *t2* can be on a uniprocessor, a tightly-coupled or a distributed multiprocessor.

(6/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Ada Array Processing

How can we rewrite what's below to complete more quickly?

```
procedure sumArray is
  N: constant integer := 1000000;
  type RealArray is array(1..N) of float;
  anArray: RealArray;

  function sum(a: RealArray; first, last: integer)
    return float is
    result: float := 0.0;
  begin
    for i in first..last loop
      result := result + a(i);
    end loop;
    return result;
  end sum;

begin
  -- code to fill anArray with values omitted
  put( sum(anArray, 1, N) );
end sumArray;
```

(7/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Divide-And-Conquer via Tasks

```
procedure parallelSumArray is
  -- declarations of N, RealArray, anArray, Sum() as before ...

  task type ArraySliceAdder
    entry SumSlice(Start: in Integer; Stop: in Integer);
    entry GetSum(Result: out float);
  end ArraySliceAdder;

  task body ArraySliceAdder is
    i, j: Integer; Answer: Float;
  begin
    accept SumSlice(Start: in Integer; Stop: in Integer) do
      i:= Start; j:= Stop;           -- get inputs
    end SumSlice;

    Answer := Sum(anArray, i, j);    -- do the work

    accept GetSum(Result: out float) do
      Result := Answer;             -- report outcome
    end GetSum;
  end ArraySliceAdder;

  -- continued on next slide...
```

(8/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Divide-And-Conquer via Tasks (ii)

```
-- continued from previous slide ...

firstHalfSum, secondHalfSum: Integer;
T1, T2 : ArraySliceAdder;  -- T1, T2 start & wait on accept
begin
  -- code to fill anArray with values omitted

  T1.SumSlice(1, N/2);      -- start T1 on 1st half
  T2.SumSlice(N/2 + 1, N);  -- start T2 on 2nd half

  T1.GetSum( firstHalfSum ); -- get 1st half sum from T1
  T2.GetSum( secondHalfSum ); -- get 2nd half sum from T2

  put( firstHalfSum + secondHalfSum ); -- we're done!
end parallelSumArray;
```

Using two tasks T1 and T2, this *parallelSumArray* version requires roughly 1/2 the time required by *sumArray* (on a multiprocessor).  
Using three tasks, the time will be roughly 1/3 the time of *sumArray*.  
...

(9/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Producer-Consumer in Ada

To give the producer and consumer separate threads, we can define the behavior of one in the 'main' procedure:

and the behavior of the other in a separate task:

We can then build a Monitor-style *Buffer* task with *put()* and *get()* as (auto-synchronizing) entry procedures...

```
procedure ProducerConsumer is
  buf: Buffer;
  it: Item;

  task consumer;
  task body consumer is
    it: Item;
  begin
    loop
      buf.get(it);
      -- consume Item it
    end loop;
  end consumer;

begin -- producer task
  loop
    -- produce an Item in it
    buf.put(it);
  end loop;
end ProducerConsumer;
```

(10/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Capacity-1 Buffer

A single-value buffer is easy to build using an Ada \_\_\_\_\_:

As a *task-type*, variables of this type (e.g., *buf*) will automatically have their own thread of execution.

The body of the task is a loop that accepts calls to *put()* and *get()* in strict alternation.

This causes *buf* to alternate between being empty and nonempty.

```
task type BoundedBuffer1 is
  entry get(it: out Item);
  entry put(it: in Item);
end BoundedBuffer1;

task body BoundedBuffer1 is
  myBuffer: Item;
begin
  loop
    accept put(it: in Item) do
      myBuffer := it;
    end put;

    accept get(it: out Item) do
      it := myBuffer;
    end get;
  end loop;
end BoundedBuffer1;
```

(11/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Capacity-N Buffer

An N-value buffer is a bit more work:

We can accept any call to *get()* so long as we are not empty, and any call to *put()* so long as we are not full.

Ada provides the *select-when* statement to \_\_\_\_\_, and perform it if and only if a given condition is *true*

```
-- task declaration is as before ...
task body BoundedBuffer is
  N: constant integer := 1024;
  package Buf is new Queue(N, Items);
begin
  loop
    select
      when not Buf.isFull =>
        accept put(it: in Item) do
          Buf.append(it);
        end put;
      or when not Buf.isEmpty =>
        accept get(it: out Item) do
          it := Buf.first;
          Buf.delete;
        end get;
    end select;
  end loop;
end BoundedBuffer;
```

(12/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## MPI ...

- ... is the \_\_\_\_\_
- ... is an industry-standard library for distributed-memory parallel computing in C, C++, Fortran, with 3<sup>rd</sup> party bindings for Java, Python, R, ...
- ... was designed by a large consortium in 1994:
  - 12 companies: *Cray, IBM, Intel, ...*
  - 11 national labs: *ANL, LANL, LLNL, ORNL, Sandia, ...*
  - representatives from 16 universities
- ... has “built in” support for many parallel design patterns
- ... continues to evolve (MPI 2.0 in 1997; 3.0 in 2012; ...)

(13/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Typical MPI Program Structure

```
#include <mpi.h>                                // MPI functions

int main(int argc, char** argv) {
    int id = -1, numProcesses = -1;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);

    // program body, which usually includes
    // calls to MPI_Send() and MPI_Receive()

    MPI_Finalize();
    return 0;
}
```

(14/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## The 6 MPI Basic Functions

1. **`MPI_Init(&argc, &argv);`**
  - Set up **`MPI_COMM_WORLD`**, a “communicator”  
(The set of processes that make up the distr. computation)
2. **`MPI_Comm_size(MPI_COMM_WORLD, &numProcesses);`**
  - How many of us processes are there to attack the problem?
3. **`MPI_Comm_rank(MPI_COMM_WORLD, &id);`**
  - Which of the  $n$  processes am I?

(15/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## The 6 MPI Basic Functions (Ct'd)

4. **`MPI_Send(sendAddress, numItems, itemType, destinationID, tag, communicator);`**
  - Send the item(s) at *sendAddress* to *destinationRank*
5. **`MPI_Recv(receiveBuffer, bufferSize, itemType, senderID, tag, communicator, status);`**
  - Receive up to *bufferSize* items from *senderRank*
6. **`MPI_Finalize();`**
  - Shut down distributed computation

These 6 commands are all you need to do useful work!

(16/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College





## Other Useful Functions

```
MPI_Bcast(receiveBuffer, bufferSize, itemType,  
          senderID, tag, comm, status);
```

- Broadcast *bufferSize* items from *senderID* to everyone in *comm*

```
MPI_Reduce(sendAddress, receiveBuffer, numItems,  
           itemType, combineOp, rootID,  
           tag, comm);
```

- Use *combineOp* to combine the distributed items at *sendAddress* into *receiveBuffer* at *destinationRank*

These (and many other) commands provide simple but efficient *collective communication*...

(17/23)

©Joel C. Adams. All Rights Reserved.

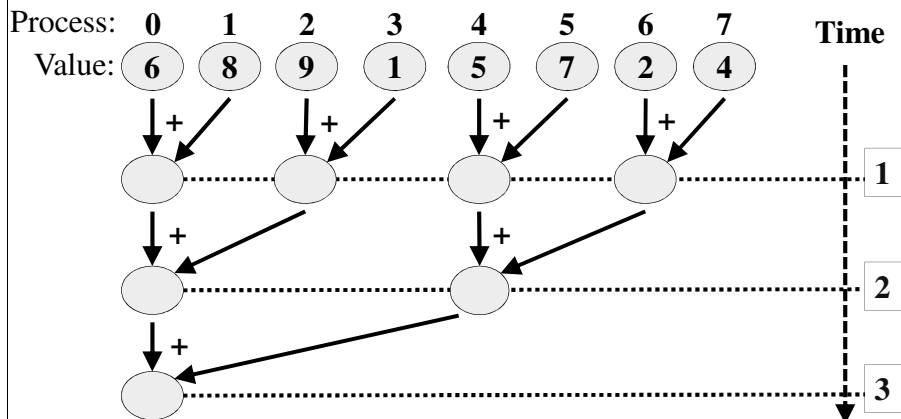
Dept of Computer Science

Calvin College



## Reduction (8 Processes)

To sum the local values of  $N = 8$  processes:



Reduction reduces the sum-time from  $O(N)$  to \_\_\_\_\_...

(18/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## A Very Simple MPI Program

```
#include <iostream>
#include <mpi.h>
using namespace std;

int main(int argc, char** argv) {
    int id = -1, n = -1;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &n);
    MPI_Comm_rank(MPI_COMM_WORLD, &id);

    int startValue = id+1;
    int square = startValue * startValue;
    int sumSquares = 0;

    MPI_Reduce(&square, &sumSquares, 1, MPI_INT,
              MPI_SUM, 0, MPI_COMM_WORLD);

    if (id == 0) {
        cout << "\nThe sum of the squares from 1 to "
              << n << " is " << sumSquares << endl;
    }

    MPI_Finalize();
    return 0;
}
```

(19/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## MPI Build and Run

To build an MPI C++ *program* from the command line:

`mpiCC program.cpp -o program`

To run an MPI *program* from the command line:

`mpirun -np N -machinefile hostFile ./program`

Launch  $N$  processes  
(each will get a unique rank)  
Vary  $N$  to test scalability

Launch those  $N$  processes  
on the computers listed in *hostFile*  
(optional ...)

Each process runs  
this same *program*  
(SPMD pattern)

(20/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Testing sumSquares

```
$ mpirun -np 1 ./sumSquares
The sum of the squares from 1 to 1 is 1

$ mpirun -np 2 ./sumSquares
The sum of the squares from 1 to 2 is 5

$ mpirun -np 3 ./sumSquares
The sum of the squares from 1 to 3 is 14

$ mpirun -np 4 ./sumSquares
The sum of the squares from 1 to 4 is 30

$ mpirun -np 128 ./sumSquares
The sum of the squares from 1 to 128 is 707264
```

(21/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Summary

- *Concurrent computations* consist of multiple entities:
  - \_\_\_\_\_ in Smalltalk, MPI    – \_\_\_\_\_ in Ada
  - \_\_\_\_\_ in C/C++, C#, Java, Go, Python, Ruby, Scala, ...
- On a shared-memory multiprocessor:
  - The \_\_\_\_\_ was the first synchronization primitive
    - *Java* provides a *Semaphore* class for synchronizing processes
  - \_\_\_\_\_ and \_\_\_\_\_ separate a semaphore's mutual-exclusion usage (locks) from its synchronization usage (c.v.s)
  - \_\_\_\_\_ are higher-level, self-synchronizing objects
    - *Java* classes have an associated (simplified) monitor
- On a distributed system:
  - *Ada* tasks provide self-synchronizing \_\_\_\_\_
  - *Erlang*, *Scala*, *MPI* support \_\_\_\_\_ between processes

(22/23)

©Joel C. Adams. All Rights Reserved.

Dept of Computer Science

Calvin College



## Summary (ii)

### Comparing Monitors and Tasks/Threads (and Coroutines):

	Has Its Own Thread	Has Its Own Execution State
Monitor		
Task/Thread		
Coroutine		

A *coroutine* (Simula, Lua) is two or more procedures that *share a single thread*, each exercising mutual control over the other:

```
procedure A;  
begin  
  -- do something  
  resume B;  
  -- do something  
  resume B;  
  -- do something  
  -- ...  
end A;
```

```
procedure B;  
begin  
  -- do something  
  resume A;  
  -- do something  
  resume A;  
  -- ...  
end B;
```

